

PRIYA YADAV

AI/ML Engineer | Generative AI & Agentic AI | LLM Evaluation

INF 524, 69120 Heidelberg, Germany | +49 01745237454 | rpsinghyadav2020@gmail.com | Portfolio Website | LinkedIn | Github

Master's student in Computational Linguistics at Heidelberg University, working as a Data Center Management student assistant at SAP SE, where I've independently built agentic AI and automation tools. Research interest in LLM evaluation and temporal reasoning in language models. Seeking full-time AI/ML, Generative AI, or Agentic AI engineering roles.

Experience

SAP SE — Working Student, Data Center Management

Aug 2025 – Present

Sankt Leon Rot, Germany

- Built an agentic skill (Python, Claude Code) deployed to SAP Joule that reasons across SQLite, ChromaDB RAG, Excel meeting minutes, Jira, and Outlook to brief management on service escalations, replacing manual review across 4 sources with a single summarized query; built weekly ingestion pipeline from Teams/OneDrive
- Built an agentic invoice reconciliation skill (Python, Claude Code) that validates build-related data against SAP's Service-Based Pricing (SBP) data, cutting a recurring monthly $\sim$ 1-hour manual check to $\sim$ 2 minutes.
- Built an agentic MRC invoice reporting skill (Python, SQLite, Claude Code) replacing manual Excel filtering with query-based rate lookup by location/data center; evolved from an Excel/VBA system with a 3-state audit trail.
- Developed a web-based SAC data preparation tool (Claude Code) featuring automated data transformation, error flagging, and inline corrections, reducing pre-upload manual prep time from 1.5–2 days to 2 hours.

Education

Ruprecht-Karls-Universität Heidelberg

Oct 2024 – Mar 2027 (expected)

Master's in Computational Linguistics · GPA: 1.8 (German scale)

Thesis (supervised by Prof. Stefan Riezler): Investigating clinical out-of-domain generalization across hospital datasets (MIMIC-IV $\rightarrow$ UMM) under distribution shift and label noise (algorithmic Sepsis-3 rules vs. expert clinician annotations)

Chandigarh University — B.E. Computer Science & Engineering

Aug 2020 – Jun 2024

81% (First Class)

Projects

Attention Mechanism Study — GLA vs. Linformer vs. Transformer (PyTorch, nanoGPT)

- Implemented and benchmarked three attention mechanisms: standard/FlashAttention, Linformer, and Gated Linear Attention (GLA), inside Karpathy's nanoGPT (124M params) on character-level Shakespeare; diagnosed Linformer's low loss as overfitting rather than genuine improvement, and traced GLA's NaN divergence to an FP16 precision failure in its cumulative log-gate term, confirming the paper's own claim that hardware-aware CUDA kernels are required.

Cold Email Generator (LangChain, ChromaDB, Streamlit)

- Built a RAG-based cold email generator that parses job descriptions into structured JSON and matches them against a portfolio-project knowledge base indexed in ChromaDB to draft targeted outreach emails.

Multimodal RAG Pipeline (LangChain, ChromaDB)

- Built a multimodal retrieval pipeline extracting structured insights from text, tables, and images in PDFs using LangChain with Llama 8B/11B, embedding document summaries for faster, more accurate retrieval.

Skills

AI/ML & Generative AI : Machine Learning, Deep Learning, Generative AI, LLMs, LLM Evaluation, Agentic AI, RAG, Prompt Engineering, Tool Orchestration, Fine-tuning

AI/ML Tools & Frameworks: Python, PyTorch, Scikit-learn, LangChain, ChromaDB, Claude Code

SAP Technologies: SAP Joule, SAP Analytics Cloud (SAC), Service-Based Pricing (SBP)

Data & Analytics: SQL, SQLite, Excel, VBA, Power BI, Tableau

Languages: English (C1, Fluent) · German (B1, Improving) · Hindi (Native)